

TAU + MPI

**Profiling and Tracing the
SUMMA Algorithm with TAU**

What and Why?

Agenda

- Define Profiling vs Tracing
- SUMMA Review
 - The algorithm
 - Data movement + costs
- TAU + Application:
 - Compile
 - Profile
 - Trace

Why?

- Performance analysis is tedious:
 - Instrumenting code by hand is repetitive
 - Adding new function-calls to code requires re-instrumenting
 - Setting up profiling and tracing code takes time
- TAU:
 - Instruments our code automatically at compile time
 - Re-written code can be re-instrumented by re-compiling
 - Puts time to instrument code back in the programmer's hands


Profiling vs. Tracing

Profiling vs. Tracing

Profiling is about how much time/resources a given program and subdivisions of said program uses. It focuses on research usage.

Tracing is about the flow of control of a program and its subdivisions: what procedures are in control at any given time, when does context switching occur, and when do system events happen across processes.

SUMMA Review



SUMMA Communication

SUMMA Algorithm

```
double* Alocal; // Column-wise random nxn matrix (local block of A)
double* Blocal; // Row-wise nxn identity matrix (local block of B)
double* Clocal; // Column wise nxn zero matrix (local block of product)
double* Atemp;  // (n/sqrt(p)) x b broadcast matrix for A
double* Btemp;  // b x (n/sqrt(p)) broadcast matrix for B
int mloc, nloc = n/sqrt(p); // size of local matrices


for (int i = 0; i < sqrt(p); i++) {
    for (int j = 0; j < m / (b*sqrt(p)); j++) {
        memcpy(Atemp, Alocal+b*j*mloc, b*mloc*sizeof(double));
        memcpy(Btemp, Blocal+b*j*nloc, b*nloc*sizeof(double));
        MPI Bcast(Atemp, b*mloc, MPI DOUBLE, i, row comm);
        MPI Bcast(Btemp, b*nloc, MPI DOUBLE, i, col_comm);
        local_gemm(Atemp, Btemp, Clocal, mloc, b);
    }
}
```

$$\gamma \cdot O\left(\frac{n^3}{p}\right) + \alpha \cdot O\left(\frac{n}{b} \log p\right) + \beta \cdot O\left(\frac{n^2}{\sqrt{p}}\right)$$

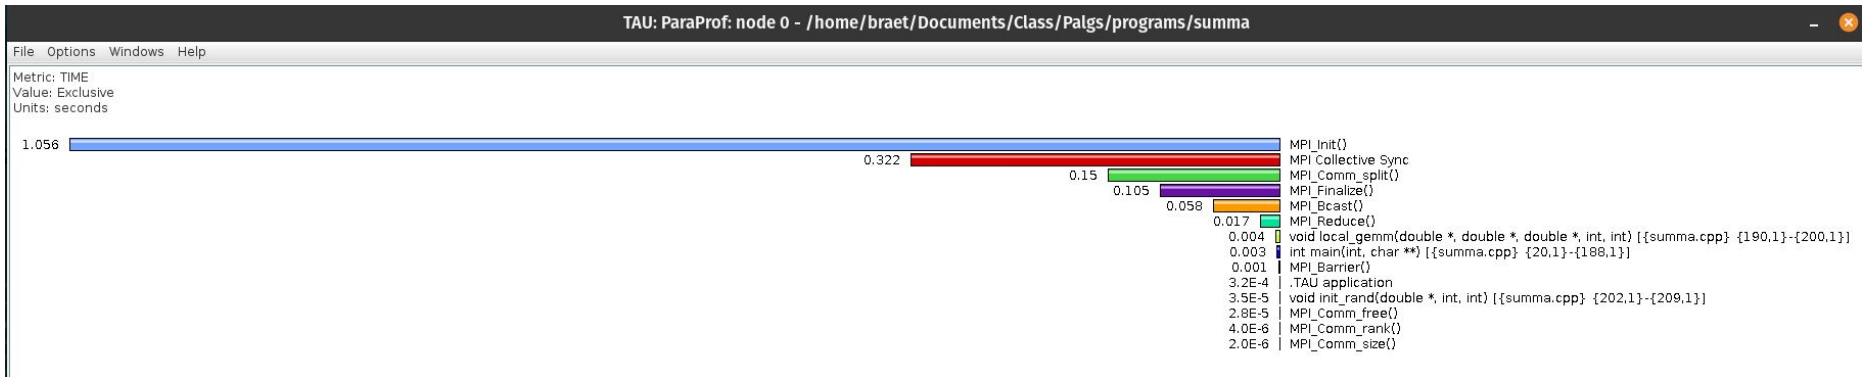
SUMMA Cost

Compile + Profile + Trace

paraprof



paraprof



pprof

FUNCTION SUMMARY (mean):

%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive usec/call	Name
100.0	0.355	3,293	1	1	3293428	.TAU application
100.0	87	3,293	1	3013	3293073	int main(int, char **)
59.7	468	1,964	2000	2000	982	MPI_Bcast()
45.7	1,505	1,505	2001	0	752	MPI Collective Sync
27.9	920	920	1	0	920400	MPI_Init()
3.3	109	109	2	0	54883	MPI_Comm_split()
3.0	98	98	1	0	98979	MPI_Finalize()
2.7	88	88	1000	0	89	void local_gemm(double *,
0.6	10	19	1	1	19782	MPI_Reduce()
0.1	2	2	1	0	2173	MPI_Barrier()
0.0	0.24	0.24	1	0	240	void init_rand(double *,
0.0	0.0434	0.0434	2	0	22	MPI_Comm_free()
0.0	0.00308	0.00308	3	0	1	MPI_Comm_rank()
0.0	0.00152	0.00152	1	0	2	MPI_Comm_size()

→ `summa git:(main)` ✕

pprof

FUNCTION SUMMARY (mean):


%Time	Exclusive msec	Inclusive total msec	#Call	#Subrs	Inclusive usec/call	Name
100.0	0.245	1,926	1	1	1926183	.TAU application
100.0	0.938	1,925	1	313	1925937	int main(int, char **)
49.8	958	958	1	0	958707	MPI_Init()
36.8	184	708	200	200	3544	MPI_Bcast()
27.3	526	526	201	0	2617	MPI_Collective Sync
5.9	113	113	2	0	56971	MPI_Comm_split()
4.5	86	86	1	0	86533	MPI_Finalize()
2.5	48	48	100	0	484	void local_gemm(double *
0.3	5	5	1	0	5952	MPI_Barrier()
0.1	0.22	2	1	1	2473	MPI_Reduce()
0.0	0.258	0.258	1	0	258	void init_rand(double *,
0.0	0.0362	0.0362	2	0	18	MPI_Comm_free()
0.0	0.00327	0.00327	3	0	1	MPI_Comm_rank()
0.0	0.00141	0.00141	1	0	1	MPI_Comm_size()

→ summa git:(main) ✕

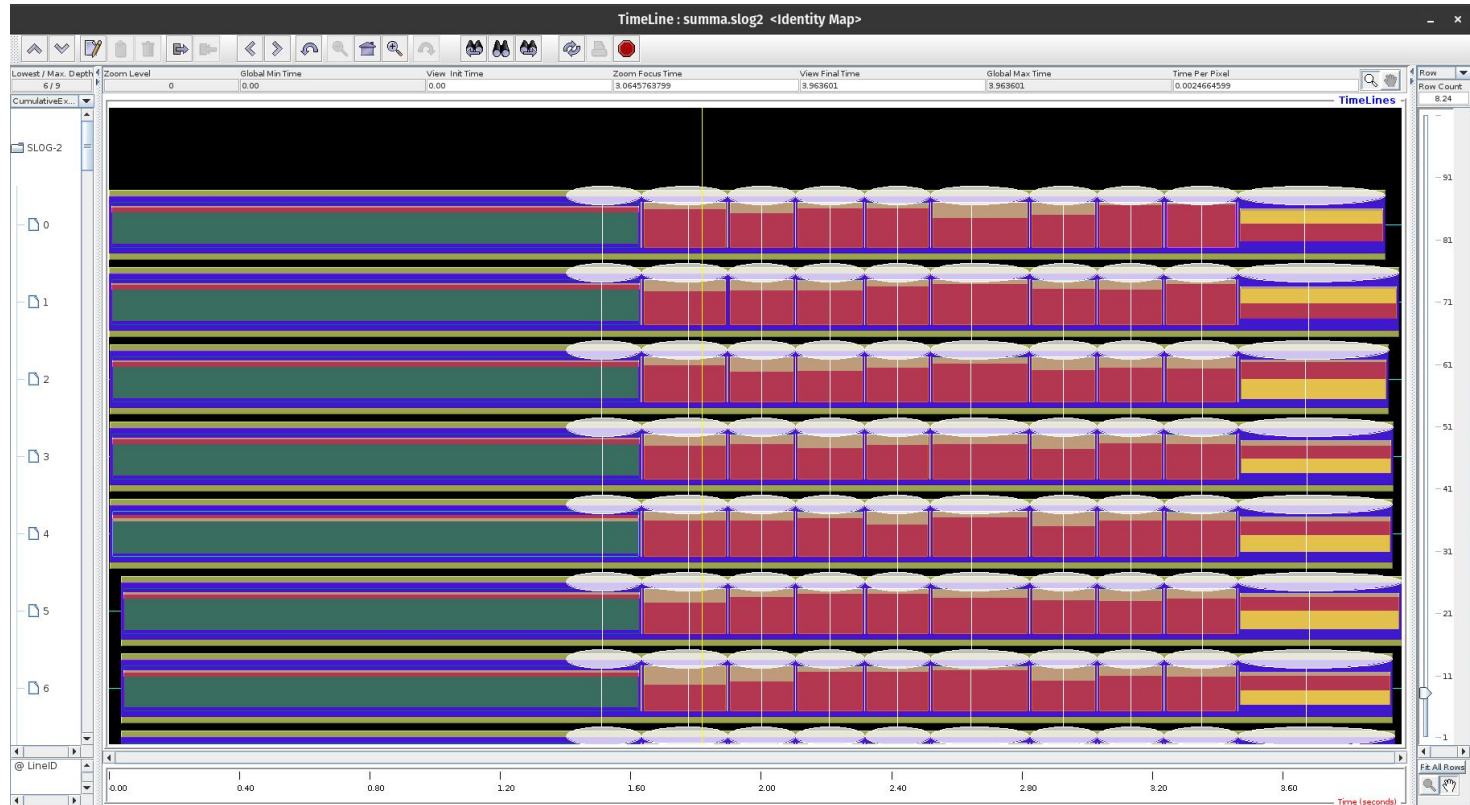
pprof

```
→ summa git:(main) X bash tau2csv.sh  
1866,1865,849,778,586,89,87,49,7,2,0.249,0.0337,0.00328,0.00139  
→ summa git:(main) X
```

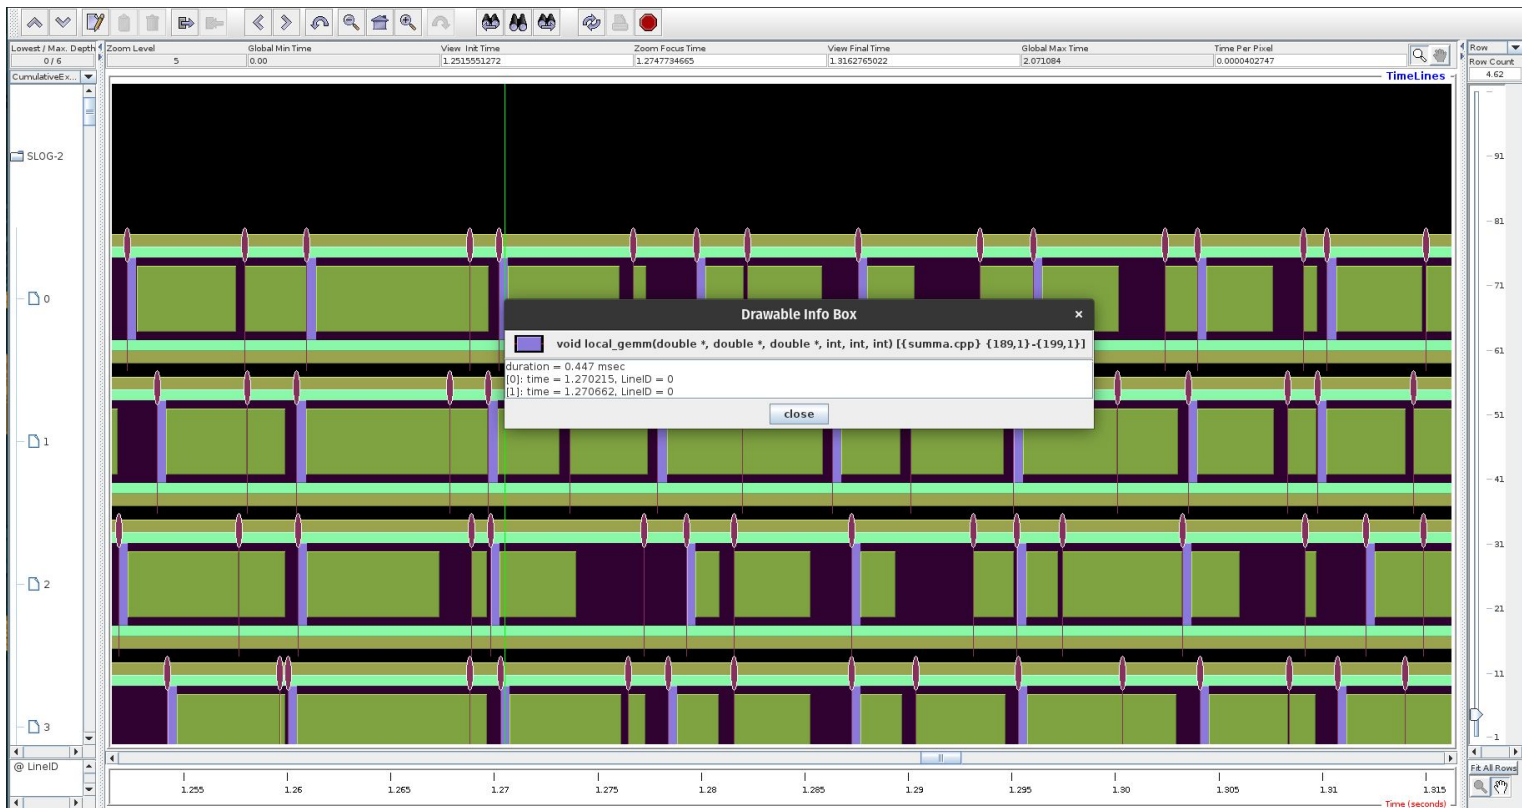

jumpshot



jumpshot



jumpshot



Q+A